

# MapReduce

CS5300

Hugo Hache, Chet Mancini, Sean Ogden



# Summary

## Challenges:


Hadoop setup  
and conventions



## Current State:

- Working locally and on Amazon EMR
- Extra credit completed
- Code includes comprehensive JUnit test cases.

# Indexing Format



## Line Number:

transform line # from input to (i, j) components

## Global:

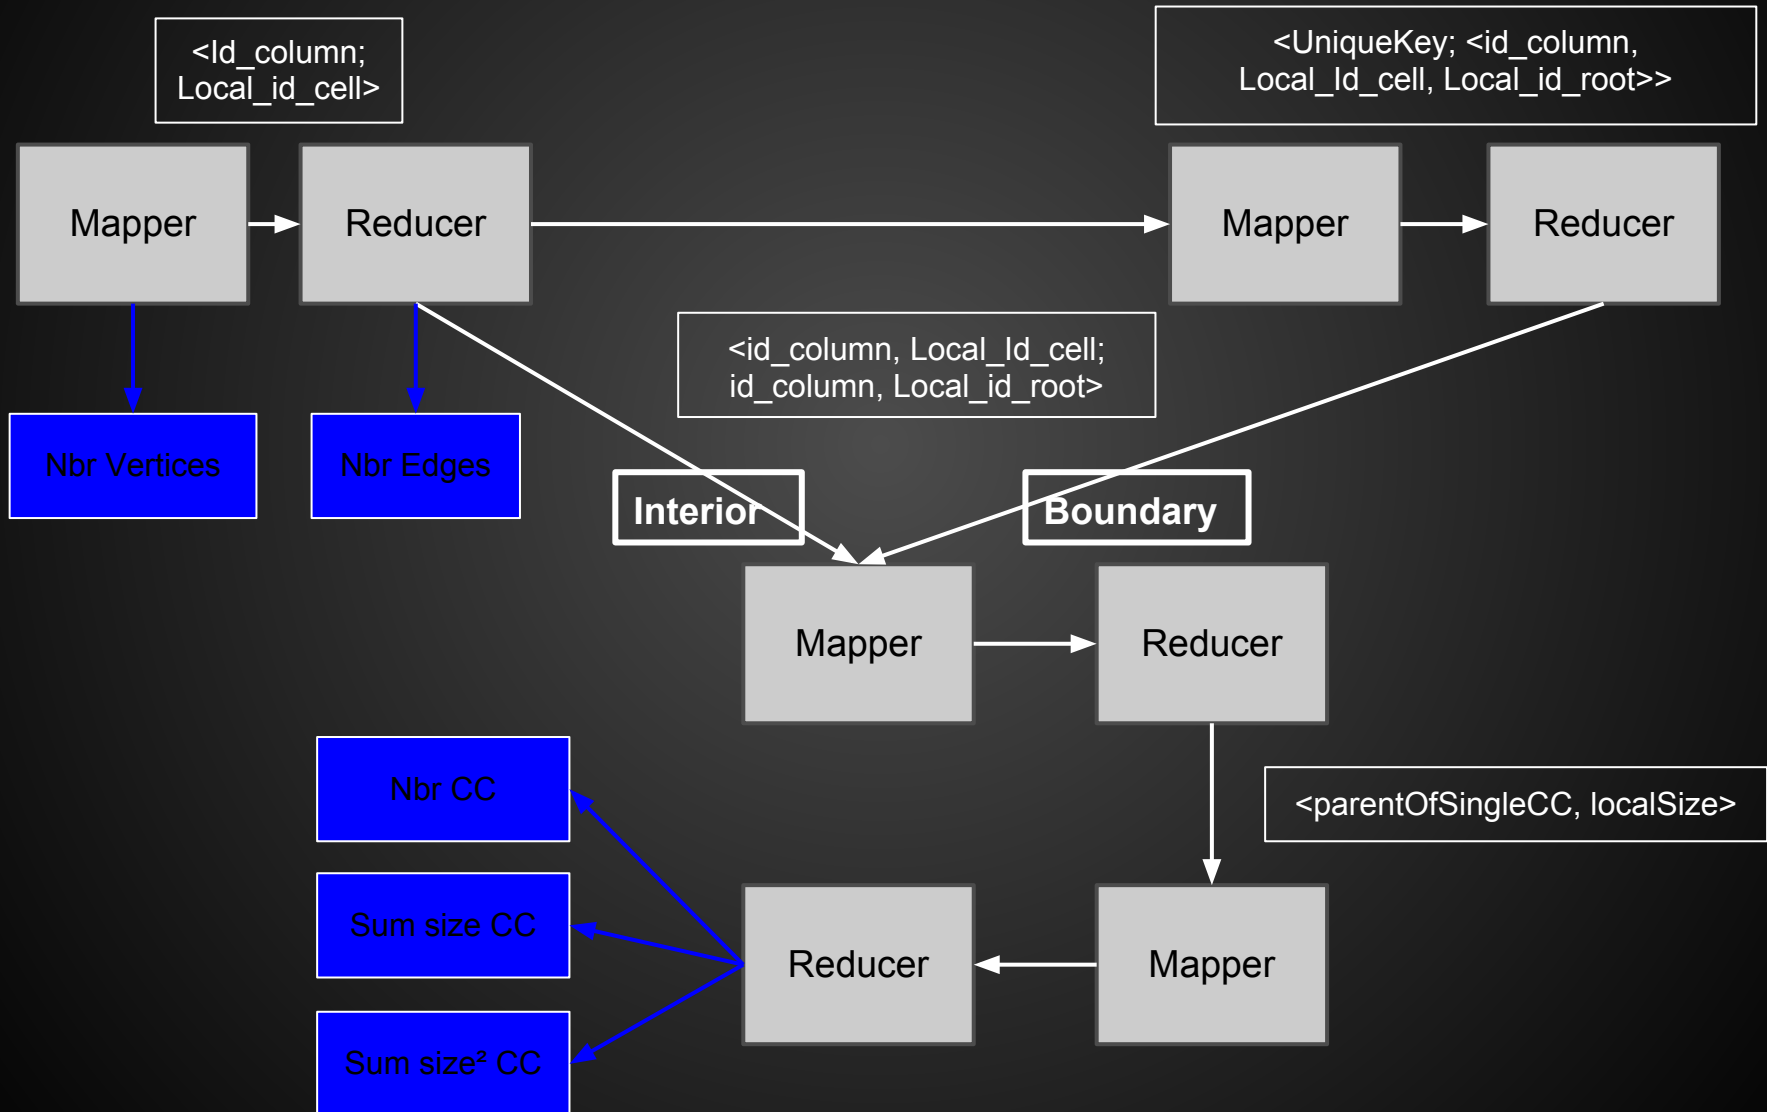
Start at southwest and count up each column.

$$\text{index} = i + (N * j)$$

## Local:

Same scheme as global, but specific to column group. Column groups also have index from 0.

# A Computation in Four Passes



# Formula of statistics

- Average Size CC

$$\text{sum size CC} / \text{nbr CC}$$

- Weighted Average Size CC =

$$\text{sum size}^2 \text{ CC} / \text{sum size CC}$$

- Average Burn Count =

$$\begin{aligned} & (\# \text{ vertices} / \# \text{ cells}) \\ & * \text{weighted avg size CC} \end{aligned}$$

# The Union Find Algorithm

- Mapped all global indices to complex numbers  $\langle G, i \rangle$
- Each column group gets  $G$ =group number,  $i = 0, \dots, N$  where  $N$  is the number of positions in a group
- Points are indexed from bottom to top, left to right.
- If two points overlap, like  $\langle 0, 30 \rangle$  and  $\langle 1, 0 \rangle$ , then the first one is strictly less than the second, so standard union find works without modification in the second pass

# Example

First phase: partitioning

<0,4>	<0,9>	<0,14>/<1,4>	<1,9>	<1,14>
<0,3>	<0,8>	<0,13>/<1,3>	<1,8>	<1,13>
<0,2>	<0,7>	<0,12>/<1,2>	<1,7>	<1,12>
<0,1>	<0,6>	<0,11>/<1,1>	<1,6>	<1,11>
<0,0>	<0,5>	<0,10>/<1,0>	<1,5>	<1,10>

# Example

Second phase: boundary columns only

$\langle 0,4 \rangle$	$\langle 0,14 \rangle$	$\langle 1,4 \rangle$	$\langle 1,14 \rangle$
$\langle 0,3 \rangle$	$\langle 0,13 \rangle$	$\langle 1,3 \rangle$	$\langle 1,13 \rangle$
$\langle 0,2 \rangle$	$\langle 0,12 \rangle$	$\langle 1,2 \rangle$	$\langle 1,12 \rangle$
$\langle 0,1 \rangle$	$\langle 0,11 \rangle$	$\langle 1,1 \rangle$	$\langle 1,11 \rangle$
$\langle 0,0 \rangle$	$\langle 0,10 \rangle$	$\langle 1,0 \rangle$	$\langle 1,10 \rangle$

- $0,10$  is to the right of  $\langle 1,0 \rangle$ , so algorithm will catch if root is lower in group 0 for that vertex



# Example

Third phase: partitioning with adjusted boundaries

<0,4>	<0,9>	<1,4>	<1,9>	<1,14>
<0,3>	<0,8>	<1,3>	<1,8>	<1,13>
<0,2>	<0,7>	<1,2>	<1,7>	<1,12>
<0,1>	<0,6>	<1,1>	<1,6>	<1,11>
<0,0>	<0,5>	<1,0>	<1,5>	<1,10>

- When there is overlap, choose the higher value (for example, <1,0> instead of <0,10>)

# Union Find: 1st Pass

- Start from bottom left point, working up to top left
  - at each vertex, check if there is another vertex below or to the left of it
  - if so, merge their roots by walking up the tree until the parent of a vertex is itself
  - replace the root that has the higher number with the lower numbered root
- Proceed to next column and work bottom to top
- continue.

# Union Find: 2nd Pass

- Start at the bottom left again
- repeat all steps from the first pass
- lowest index from every component will definitely be encountered before any other, so two passes are sufficient

# Our Results (non-diagonal)

From NetID "974"

Our statistics on productions.txt with size 10000\*10000:

<b># of vertices</b>	58997294
<b># of edges</b>	69608194
<b># CC</b>	2897988
<b>Avg CC Size</b>	20.358019
<b>Weighted Avg CC Size</b>	7543.9851
<b>Burn Count</b>	4450.74707628

# Performance

# Instances	Minutes
10	25
20	19
30	14

## Results w/ Diagonal Added

